

Provider & Practice Directory Update Pipeline

HealthLynked challenge — Option C: working prototype + production architecture.

Executive summary

Healthcare provider data decays constantly — providers move, practices merge or rebrand, phones and suites change, clinicians retire. Maintaining a directory by hand is expensive and doesn't scale. I built a **repeatable, cost-efficient pipeline** that keeps a provider **and** practice directory accurate, and I can **measure** its accuracy rather than assert it.

The thesis in one line: most provider and practice updates already sit in free, authoritative U.S. government datasets (NPPES + CMS), so the optimal pipeline is a **deterministic-first funnel** that resolves the bulk of records at ≈\$0 marginal cost and reserves web scraping, LLMs, and human reviewers for the small residual where structured sources are silent or in conflict.

Three things set this submission apart:

1. **It runs.** A one-command prototype reconciles the brief's exact `HL_001` example, emits the brief's exact JSON schema, and exercises every evaluation criterion end-to-end — offline, deterministically, with no API keys.
 2. **Accuracy is measured.** Because the government sources are versioned monthly, I replay a past snapshot against a later one and report real precision, recall, and a calibration curve. A representative 5,000-record backtest shows **≈98% recall, ≈97% auto-update precision, 100% routing accuracy** (Brier ≈ 0.05).
 3. **The timing is a legal mandate.** The **REAL Health Providers Act** (signed 2026-02-03) will require Medicare Advantage plans to verify every directory field every 90 days, remove departed providers within 5 days, and **publicly report an accuracy score** from PY2028. A pipeline that emits a defensible, auditable accuracy score is the compliance product HealthLynked's payer customers will be required to buy.
-


1. What the evaluation rewards, and the design choice that earns it

Criterion	Design choice
Accuracy	Anchor on authoritative structured sources (NPPES/CMS) before any LLM/scrape; never auto-update on a single weak source — and prove it with a historical-snapshot backtest.
Scalability	National bulk files, embarrassingly parallel by state/ZIP; blocking keeps matching sub-quadratic; Splink-ready for 100M+.
Cost efficiency	A free deterministic funnel resolves the large majority; LLM/scrape/human touch only the residual. Flat cost in volume.
Practicality	Free OSS + free government APIs; no exotic infra; a 2–3 person team can run it.
Explainability	Every recommendation carries the scoring math and the source snapshots that drove it.
Data quality	Canonical normalization for address (+ Census geocode), phone (E.164), specialty (NUCC taxonomy), and name.
Source reliability	Explicit per-source, per-field reliability weights; independence-aware corroboration; conflict baked into the score.
Human-review design	Risk-tiered routing so only true conflicts/low-confidence reach a human; a working review dashboard.
Audit trail	Append-only, content-hashed event store; every change links to source evidence and the confidence computation.

2. Architecture

A deterministic-first funnel routes each record only as far down a **cost/authority gradient** as it needs to go. The expensive tiers (LLM extraction, human review) only ever see records the free tiers couldn't resolve — that single choice is what keeps the cost curve flat as volume grows.

Provider & Practice Directory Update Pipeline — Agent Workflow



Naming the boxes "agents" is fair, but most are plain deterministic code: only triage, the gated harvest, and LLM extraction are agentic; normalization, matching, scoring, and the audit store are ordinary, testable functions. That honesty is a feature — it is what makes the system cheap and reproducible.

3. Data sources & reliability (free first)

Source	Authoritative for	Access	Cost
NPPES NPI Registry	Identity, NPI, taxonomy→specialty, practice/mailing address, phone, deactivation. Type-1 = providers, Type-2 = practices (both halves of the brief).	Free API (no key) + monthly bulk	\$0
CMS Doctors & Clinicians (mj5m-pzi6)	Group/practice affiliation (PAC ID), practice locations, active enrollment.	Free API + bulk	\$0
Historical NPPES/CMS snapshots (NBER mirror)	Ground truth for the backtest — what really changed month to month.	Free	\$0
US Census Geocoder	Address standardization/validation + coordinates.	Free API (no key)	\$0
State medical boards	License status (active/inactive/disciplinary).	Per-state scrape	scrape only
Practice website	Current phone/address/roster — strong corroborator.	Scrape + LLM (Tier 3 only)	\$ (residual only)

Paid vendors (Google Places, Melissa, commercial provider data) are deliberately **avoided** on the default path — they are the easy way to blow the cost budget the brief explicitly penalizes.

4. Confidence scoring (the formula)

For a field `f` and candidate value `v` asserted by sources `S`:

$$\text{field_conf}(f, v) = R(v) \cdot A(v)$$
$$A(v) = \frac{\sum_{\{s: \text{val}=v\}} w(s, f) \cdot \phi(s)}{\sum_{\{s \in S\}} w(s, f) \cdot \phi(s)} \quad (\text{agreement share})$$
$$R(v) = 1 - \prod_{\{s: \text{val}=v\}} (1 - w(s, f) \cdot \phi(s)) \quad (\text{noisy-OR strength})$$

- `w(s, f)` is a **per-field reliability weight** for the source class; `φ(s)` is a **freshness** factor that decays stale source data.
- `A` drags confidence toward the middle when sources disagree (genuine conflict → ambiguous score → human review). `R` rewards multiple independent confirmations.
- Independence is enforced.** NPPES and CMS are *not* independent — both are ultimately self-reported by the provider to CMS — so same-class sources are de-correlated (the 2nd+ source in a class contributes only `1-p` of its weight), and the corroboration gate requires agreement across **≥2 distinct independence classes**: provider-self-reported gov (NPPES), claims/enrollment gov (CMS), practice web, and regulatory board.

The brief's conflict example falls straight out: the practice website and CMS show a new address while NPPES still shows the old one, so the address confidence collapses to ~0.58 and the record is routed to a human — while the phone, agreed by two sources, scores 0.88.

5. Decision & safe-update rules

Field **risk classes** govern thresholds and corroboration:

Class	Fields	Rule
Low	phone, suite, website	auto-update at <code>conf ≥ 0.85</code> , ≥1 class
Medium	full address (move), specialty, practice name, affiliation	auto-update at <code>conf ≥ 0.90</code> and ≥2 independent classes
High	provider name, NPI , active status, practice merge/close	never silent — always routed to a human

```
if no field differs                → NO_CHANGE
elif a high-risk source contradicts current → HUMAN_REVIEW (dissent safety net)
elif any changed field is in conflict → HUMAN_REVIEW
elif a high-risk field changed      → HUMAN_REVIEW
elif min(conf) ≥ threshold and corroborated → AUTO_UPDATE
elif min(conf) ≥ τ_low              → HUMAN_REVIEW
else                               → DISCARD / HOLD
```

Two safety choices worth calling out:

- **NPI never auto-updates.** An NPI that maps to a different entity is a data-entry error or a merged duplicate → a repair/dedup ticket, never a silent field write.
- **Downgrades are never silently ignored.** If a credible source reports a provider *inactive* while the record (backed by a stale stronger source) still shows *active*, the record is routed to review even though nothing "changed" — missing a deactivation is exactly the trust-eroding error the brief cares about.

6. Normalization, matching & lifecycle signals

- **Normalize:** address → canonical components + Census geocode (suite handled separately so a suite change never looks like a move); phone → E.164; specialty → NUCC taxonomy code; name → structured parts with credentials pulled out.
- **Match (Fellegi–Sunter, Splink-ready):** dedup by blocking on NPI and `(soundex(name), ZIP)`; **practice-location matching uses geocoded proximity**, not fragile string similarity.
- **Lifecycle signals:** the pipeline labels *provider movement*, *practice relocation*, *closure* (deactivation), *rebrand*, and *NPI repair* as structured events — not guesses.

7. Triage — the cost lever

The cheapest verification is the one you skip. Each record is risk-scored so a cycle actively harvests only the top of the queue; everything else rides the free monthly batch diff.

```
risk = w1·staleness(days vs 90-day cadence) + w2·field_volatility
      + w3·hard_signal(deactivation / diff / bounce) + w4·directory_importance
```

8. Validation & calibration — *measured, not asserted*

Because the authoritative sources are versioned monthly, I replay snapshot **T** as "the directory" and score the pipeline's proposals against snapshot **T+Δ** as ground truth — labels for free, at national scale. The same harness is a regression test and produces exactly the accuracy score the REAL Act now mandates.

The repo ships three fidelity levels behind one metric implementation: a fast offline stand-in (CI/regression), a **synthetic two-snapshot world at scale** (with a realistic fraction of stale/wrong sources), and a builder that diffs two real NBER NPPES snapshots. A representative `--scale 5000` run:

Metric	Result
Recall of real changes	97.7%
Auto-update precision (1 – wrong silent writes)	96.8%
Routing accuracy	100%
Brier score / ECE	0.05 / 0.08 (well-calibrated)

Confidence is then calibrated (isotonic/Platt) so a published `0.90` means ~90% of such updates are correct — what makes the audit trail and the mandated accuracy score defensible to a payer's compliance auditor.

One honest limit: a snapshot backtest measures agreement with the *authoritative record*, and NPPES contact data is self-reported and can lag the real world. So for the fields whose ground truth lives off the government grid — a working phone, a current suite — production also holds out a small **hand-labeled gold set** (a few hundred records verified by phone/web) to measure on-the-ground accuracy and keep the Tier-3 enrichment honest. The backtest covers identity / taxonomy / affiliation / deactivation at national scale; the gold set covers contact accuracy in depth.

9. Cost model — per 1,000 records

Stage	Records / 1,000	Subtotal
Tier 0/1 batch reconcile (NPPES+CMS bulk, amortized)	1,000	~\$0.50
NPPES API + Census geocode (targeted, free)	~300	\$0
Tier 3 web fetch + LLM extraction (residual only)	~100	\$1–3
Human review (true conflicts only)	~30–50	\$30–50
Total		≈ \$32–54 / 1,000 ≈ \$0.03–0.05 per record

Cost is **dominated by human review** — which the entire architecture exists to minimize — and is **flat in volume** because Tiers 0–1 are free batch. A naive "LLM-every-record + paid address API" approach lands 10–20× higher and is less accurate.

10. Scalability

The national NPPES file is $\approx 8\text{--}9\text{M}$ NPIs (~ 1 GB) and full reconciliation is embarrassingly parallel (partition by state/ZIP). The per-record pipeline is a pure map — drop it behind a worker pool or queue. Matching blocks so it never goes quadratic and swaps to Splink (100M+ on Spark). Live calls retry with backoff and cache on disk, so re-runs and the 90-day re-verify cadence don't re-hit the network.

11. Human review & audit trail

A lean Streamlit dashboard shows one card per record needing a human — each proposed change side-by-side with its confidence, supporting sources, and conflict flag, with the queue sorted so the highest-risk and conflicting records surface first. Every approve/reject is logged as a **training label** that tunes thresholds over time, so the system grows more autonomous as it earns trust.



Provider Directory — Human Review Queue

Only the records the pipeline could not (or should not) auto-update reach this screen. Every approve / reject is logged as a label that tunes the scoring thresholds over time, so the system grows more autonomous as it earns trust.

3

In review queue

3

Pending your review

0

Reviewed this session

HL_004 NPI 1555666779 **HIGH RISK**

High-risk field(s) changed (active); never updated silently - routed to human confirmation.

Active

Active → Deactivated

NPI Registry

State Medical Board

0.99

✓ Approve

✗ Reject

Snooze

HL_001 NPI 1234567890 **MEDIUM RISK** **△ SOURCE CONFLICT**

NPI 1234567890 fails its check digit; proceeding but flagged for repair. Practice Website, CMS Doctors & Clinicians and NPI Registry disagree on address; manual verification recommended.

Address

~~100 Main St, Naples, FL 34102~~ → **250 Health Park Dr, Fort Myers, FL 33908**

CMS Doctors & Clinicians

Practice Website

△ conflict

0.58

Phone

~~239-555-1234~~ → **239-555-9000**

NPI Registry

Practice Website

0.88

✓ Approve

✗ Reject

Snooze

HL_003 NPI 1234567890 **MEDIUM RISK** **△ SOURCE CONFLICT**

NPI 1234567890 fails its check digit; proceeding but flagged for repair. Practice Website, CMS Doctors & Clinicians and NPI Registry disagree on address; manual verification recommended.

Address

~~100 Main Street, Suite 2, Naples, FL 34102~~ → **250 Health Park Dr, Fort Myers, FL 33908**

CMS Doctors & Clinicians

Practice Website

△ conflict

0.58

Phone

~~239-555-1234~~ → **239-555-9000**

NPI Registry

Practice Website

0.88

✓ Approve

✗ Reject

Snooze

The review queue on the demo data: the high-risk deactivation (HL_004) is surfaced first, then the two address conflicts where the practice website and CMS disagree with NPPES — each field showing its calibrated confidence, supporting sources, and conflict flag. A reviewer approves, rejects, or snoozes per record; the green/amber/red confidence bars make the safe calls obvious at a glance.

Every proposed or applied change writes one immutable, **content-hashed** audit event: old→new value, the source snapshots (with weights and freshness), the score breakdown, the decision, the actor, and the pipeline version — so any update is reproducible and defensible months later.

12. Roadmap (also the 3-month consulting plan)

Phase	Weeks	Deliverable
0 — MVP (this submission)	—	Runnable funnel, schema-compliant output, audit log, dashboard, backtest.
1 — Structured backbone	1–4	Bulk NPPES+CMS ingestion; full-directory monthly reconcile; backtest harness on real snapshots from day one.
2 — Matching & triage	4–8	Splink dedup/movement in production; risk-based triage; weekly incrementals + deactivation diffs.
3 — Residual enrichment	8–12	Gated Tier-3 agent (web/board) with budget caps + caching; dashboard live; threshold tuning from reviewer labels.

KPIs, measured continuously by the backtest: auto-update precision $\geq 99\%$ on the production threshold, human-review $\leq 5\%$, cost $\leq \$0.05/\text{record}$, freshness within the 90-day window, plus a published, calibrated accuracy score.

13. Risks & mitigations

Risk	Mitigation
NPPES address is self-reported and can lag	Treat as strong-but-not-sole; require corroboration for moves; Tier-3 confirms.
State-board scraping is brittle (50 sites)	Build incrementally, highest-volume states first; board is a corroborator, not a hard dependency.
LLM hallucination	LLM only <i>extracts</i> from fetched text; the value must still clear the scoring gate; snapshot stored.
Over-aggressive auto-update	Conservative thresholds, min-field-confidence decision, high-risk always human-confirmed, full reversibility via the audit log.

Appendix — the prototype

```
pip install -e ".[dev,calibration]"
directory-pipeline demo           # end-to-end offline showcase
directory-pipeline backtest --scale 5000  # the measured accuracy curve
streamlit run dashboard/app.py      # human-review dashboard
```

The demo reconciles a 6-record sample (including the brief's HL_001 and a Type-2 practice record), prints triage, recommendations, lifecycle signals, duplicate detection, and a measured backtest — all logged to a content-hashed audit store.

=====

Provider & Practice Directory Update Pipeline - offline demo

=====

6 sample records | sources: NPPEs + CMS + web/board (fixtures)

Verify queue (triage -- most at-risk first):

[0.63] HL_004 (1980d since last verify (> 90d window))
 [0.51] HL_002 (1490d since last verify (> 90d window))
 [0.48] HL_006 (1320d since last verify (> 90d window))
 [0.41] HL_001 (1016d since last verify (> 90d window))
 [0.41] HL_003 (1016d since last verify (> 90d window))
 [0.19] HL_005 (104d since last verify (> 90d window))

[?] HL_001 John Smith, MD

action : human_review (overall 0.73)
 reason : NPI 1234567890 fails its check digit; proceeding but flagged for repair. Practice Website, CMS Doctors & Clinicians and NPI...
 change : address: '100 Main St, Naples, FL 34102' -> '250 Health Park Dr, Fort Myers, FL 33908' (conf 0.58; Practice Website, CMS D...
 change : phone: '239-555-1234' -> '239-555-9000' (conf 0.88; NPI Registry, Practice Website)
 signal : provider movement: 100 Main St, Naples, FL 34102 -> 250 Health Park Dr, Fort Myers, FL 33908
 geocode : address validated (US Census)

[+] HL_002 Jane A Doe, NP

action : auto_update (overall 0.96)
 reason : NPI 1002003006 recovered by name/org search (agreement 1.00). Updated address, phone confirmed by independent reliable sour...
 change : address: '300 1st St, Fort Myers, FL 33901' -> '501 HEALTHPARK BLVD, FORT MYERS, FL 33908' (conf 0.96; Practice Website, N...
 change : phone: '239-555-2000' -> '239-555-7777' (conf 0.96; NPI Registry, Practice Website, CMS Doctors & Clinicians)
 signal : provider movement: 300 1st St, Fort Myers, FL 33901 -> 501 HEALTHPARK BLVD, FORT MYERS, FL 33908
 geocode : address validated (US Census)

[?] HL_003 John Smith M.D.

action : human_review (overall 0.73)
 reason : NPI 1234567890 fails its check digit; proceeding but flagged for repair. Practice Website, CMS Doctors & Clinicians and NPI...
 change : address: '100 Main Street, Suite 2, Naples, FL 34102' -> '250 Health Park Dr, Fort Myers, FL 33908' (conf 0.58; Practice W...
 change : phone: '239-555-1234' -> '239-555-9000' (conf 0.88; NPI Registry, Practice Website)
 signal : provider movement: 100 Main Street, Suite 2, Naples, FL 34102 -> 250 Health Park Dr, Fort Myers, FL 33908
 geocode : address validated (US Census)

[?] HL_004 Robert Lee, MD

action : human_review (overall 0.99)
 reason : High-risk field(s) changed (active); never updated silently - routed to human confirmation.
 change : active: 'Active' -> 'Deactivated' (conf 0.99; NPI Registry, State Medical Board)
 signal : provider inactive/retired
 geocode : address validated (US Census)

[=] HL_005 Maria Garcia, MD

action : no_change (overall 0.89)
 reason : All checked fields match authoritative sources; record confirmed accurate.
 geocode : address validated (US Census)

[+] HL_006 Gulf Coast Women's Health

action : auto_update (overall 0.94)
 reason : Updated address, phone confirmed by independent reliable sources (min confidence 0.94 >= 0.90).
 change : address: '850 6th Ave N, Naples, FL 34102' -> '3300 PINE RIDGE RD, NAPLES, FL 34109' (conf 0.95; Practice Website, NPI Reg...
 change : phone: '239-555-6000' -> '239-555-6100' (conf 0.94; Practice Website, NPI Registry, CMS Doctors & Clinicians)
 signal : practice relocation: 850 6th Ave N, Naples, FL 34102 -> 3300 PINE RIDGE RD, NAPLES, FL 34109
 geocode : address validated (US Census)

Duplicate clusters detected:

[DUP_001] HL_001, HL_003 (score 1.00; identical NPI 1234567890)

Audit log written to: audit_log.jsonl (10 events)

Records awaiting human review: 3 (5 field-level events)

Backtest report (historical-snapshot replay)

=====

records evaluated : 6
 true field changes : 9
 detected (correct/all) : 9/9

recall : 100.0%
 detection precision : 100.0%
 AUTO-UPDATE precision : 100.0% (4/4)
 routing accuracy : 100.0% (6/6)

Brier score (lower=better): 0.0434

```
expected calib. error      : 0.1422
```

```
NOTE: only 9 scored changes – too few for a stable  
calibration curve. Run `directory-pipeline backtest --scale 5000`  
for the statistically meaningful curve (isotonic in production).
```

```
-----  
Launch the human-review dashboard with: streamlit run dashboard/app.py
```

Solo submission. All code is free/OSS; all data sources are free and authoritative.